

Predicting Soccer Match Outcomes

Theyab Alkhoori, Jasmine Moskowitz, Max Treusein

2025-12-02

Introduction

Around the world, soccer is a popular sport. Millions of people enjoy using each team's data to predict match outcomes. Our project aims to predict match results of the 2017 season. We will predict the 2017 season using the statistics from the 2008 to 2016 seasons, then compare our predictions with the actual outcomes. In the end we hope to know if the home team wins, ties, or loses and well as the number of goals scored during the game. Once complete we can then compare our results with the actual results of the game.

References

We plan to execute our desired tasks by using the **European Soccer Database** from Kaggle (<https://www.kaggle.com/datasets/hugomathien/soccer/data>). This source includes information from more than 25,000 games, more than 10,000 players, and eleven European nations. The data contains match statistics, results, and other information all contained in csv files. We will split the data into two parts: train and test. The training data will contain data from the season 2008 to 2016. The test data will contain data from the 2017 season.

Methods

Task 1: Data Preparation

Our first task is to clean out the data from the csv files. We want to extract all the necessary data to make it easier to complete our tasks.

Task 2: Classification

Our second task is using classification methods to predict match outcomes (home win, tie, home loss)

Task 3: Regression

Our third task is to predict the number of goals scored in a match.

Results

1. If the home team wins, ties, or loses.
2. The number of goals scored during a game.

=====

We used the **European Soccer Database** from Kaggle, which includes 11 leagues, 10,000 players, and more than 25,000.

```

library(tidymodels)

## -- Attaching packages ----- tidymodels 1.3.0 --
## v broom      1.0.8   v recipes     1.3.1
## v dials      1.4.0   v rsample     1.3.0
## v dplyr      1.1.4   v tibble      3.3.0
## v ggplot2    4.0.0   v tidyverse   1.3.1
## v infer      1.0.9   v tune        1.3.0
## v modeldata   1.4.0   v workflows   1.2.0
## v parsnip     1.3.2   v workflowsets 1.1.1
## v purrr      1.1.0   v yardstick   1.3.2

## -- Conflicts ----- tidymodels_conflicts() --
## x purrr::discard() masks scales::discard()
## x dplyr::filter()  masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## x recipes::step() masks stats::step()

library(discrim)

##
## Attaching package: 'discrim'

## The following object is masked from 'package:dials':
## 
##   smoothness

library(ISLR)
library(ggplot2)
library(dplyr)
library(readr)

##
## Attaching package: 'readr'

## The following object is masked from 'package:yardstick':
## 
##   spec

## The following object is masked from 'package:scales':
## 
##   col_factor

library(lubridate)

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
## 
##   date, intersect, setdiff, union

set.seed(445)

## Load raw data ----
match <- read.csv("Match.csv")
team <- read.csv("Team.csv")
team_attr <- read.csv("Team_Attributes.csv")

```

```

## Team-season attributes ----
team_attr |>
  mutate(
    date = as.Date(date),
    season = year(date)
  ) |>
  filter(season >= 2010, season <= 2015) |>
  select(-id, -date) |>
  group_by(team_api_id, season) |>
  summarise(
    across(where(is.numeric), ~ mean(.x, na.rm = TRUE)),
    .groups = "drop"
  ) -> team_season_attr

team |>
  select(team_api_id, team_long_name, team_short_name) |>
  right_join(team_season_attr, by = "team_api_id") -> team_season_attr

num_cols <- sapply(team_season_attr, is.numeric)

team_season_attr[num_cols] <- lapply(team_season_attr[num_cols], function(x) {
  if (all(is.na(x))) return(x)
  x[is.na(x)] <- mean(x, na.rm = TRUE)
  x
})

## Build match-level outcome and season ----
match |>
  mutate(
    date = as.Date(date),
    season = year(date),
    home_result = case_when(
      home_team_goal > away_team_goal ~ "Win",
      home_team_goal == away_team_goal ~ "Draw",
      home_team_goal < away_team_goal ~ "Loss"
    ),
    home_result = factor(home_result, levels = c("Win", "Draw", "Loss"))
  ) |>
  filter(season >= 2010, season <= 2015) |>
  select(
    match_api_id,
    season,
    date,
    home_team_api_id,
    away_team_api_id,
    home_team_goal,
    away_team_goal,
    home_result
  ) -> matches_base

## Home-team attributes ----
team_season_attr |>
  rename(

```

```

    home_team_api_id = team_api_id
) |>
rename_with(
  ~ paste0("home_", .x),
  .cols = -c(home_team_api_id, season)
) -> home_attr

## Away-team attributes ----
team_season_attr |>
  rename(
    away_team_api_id = team_api_id
) |>
  rename_with(
    ~ paste0("away_", .x),
    .cols = -c(away_team_api_id, season)
) -> away_attr

## Join attributes onto matches ----
matches_base |>
  left_join(home_attr, by = c("season", "home_team_api_id")) |>
  left_join(away_attr, by = c("season", "away_team_api_id")) -> soccer_matches

## Final train / test splits ----
train_matches <- soccer_matches |> filter(season <= 2014)
test_matches <- soccer_matches |> filter(season == 2015)

head(train_matches)

##   match_api_id season      date home_team_api_id away_team_api_id
## 1       665626  2010 2010-02-03          8635          8342
## 2       665630  2010 2010-02-04          9986          9985
## 3       665634  2010 2010-02-02          8203          9993
## 4       665665  2010 2010-01-30          8342         10001
## 5       665666  2010 2010-01-17          9985          8635
## 6       665667  2010 2010-01-16          9997          9994
##   home_team_goal away_team_goal home_result home_team_long_name
## 1             3             2        Win      RSC Anderlecht
## 2             2             3       Loss    Sporting Charleroi
## 3             1             0        Win       KV Mechelen
## 4             2             1        Win     Club Brugge KV
## 5             0             4       Loss  Standard de Liège
## 6             2             1        Win  Sint-Truidense VV
##   home_team_short_name home_team_fifa_api_id home_buildUpPlaySpeed
## 1                 AND                  229                  50
## 2                 CHA                  670                  40
## 3                 MEC                110724                  65
## 4                 CLB                  231                  35
## 5                 STL                  232                  53
## 6                 STT                  680                  50
##   home_buildUpPlayDribbling home_buildUpPlayPassing home_chanceCreationPassing
## 1           48.60451                  35                  70
## 2           48.60451                  50                  45
## 3           48.60451                  60                  50
## 4           48.60451                  40                  45

```

```

## 5          48.60451           40           55
## 6          48.60451           50           50
##   home_chanceCreationCrossing home_chanceCreationShooting home_defencePressure
## 1          50                 60           70
## 2          43                 60           70
## 3          40                 50           60
## 4          50                 45           60
## 5          55                 65           70
## 6          50                 50           50
##   home_defenceAggression home_defenceTeamWidth away_team_long_name
## 1          50                 70 Club Brugge KV
## 2          70                 70 Standard de Liège
## 3          70                 60 Beerschot AC
## 4          70                 70 KVC Westerlo
## 5          70                 65 RSC Anderlecht
## 6          50                 50 Sporting Lokeren
##   away_team_short_name away_team_fifa_api_id away_buildUpPlaySpeed
## 1          CLB                231          35
## 2          STL                232          53
## 3          BAC                675          35
## 4          WES                681          45
## 5          AND                229          50
## 6          LOK                2007         60
##   away_buildUpPlayDribbling away_buildUpPlayPassing away_chanceCreationPassing
## 1          48.60451           40           45
## 2          48.60451           40           55
## 3          48.60451           35           45
## 4          48.60451           60           40
## 5          48.60451           35           70
## 6          48.60451           60           50
##   away_chanceCreationCrossing away_chanceCreationShooting away_defencePressure
## 1          50                 45           60
## 2          55                 65           70
## 3          40                 50           70
## 4          45                 55           70
## 5          50                 60           70
## 6          40                 50           65
##   away_defenceAggression away_defenceTeamWidth
## 1          70                 70
## 2          70                 65
## 3          70                 70
## 4          70                 70
## 5          50                 70
## 6          65                 70

head(test_matches)

##   match_api_id season      date home_team_api_id away_team_api_id
## 1    1717979   2015 2015-02-03        9987        8573
## 2    1717985   2015 2015-01-16       8342        8203
## 3    1717986   2015 2015-01-18       8635        9989
## 4    1717987   2015 2015-01-17       8573        8571
## 5    1717988   2015 2015-01-17       9986        9984
## 6    1717989   2015 2015-01-17      274581        9991
##   home_team_goal away_team_goal home_result home_team_long_name

```

```

## 1      1      1      Draw      KRC Genk
## 2      1      1      Draw      Club Brugge KV
## 3      3      0      Win       RSC Anderlecht
## 4      1      7      Loss      KV Oostende
## 5      0      2      Loss      Sporting Charleroi
## 6      1      3      Loss      Royal Excel Mouscron
##   home_team_short_name home_team_fifa_api_id home_buildUpPlaySpeed
## 1           GEN          673          58
## 2           CLB          231          42
## 3           AND          229          52
## 4           OOS          682          52
## 5           CHA          670          60
## 6           MOP         111560         50
##   home_buildUpPlayDribbling home_buildUpPlayPassing home_chanceCreationPassing
## 1             52            38            30
## 2             52            45            42
## 3             46            51            60
## 4             49            54            45
## 5             47            53            48
## 6             50            50            50
##   home_chanceCreationCrossing home_chanceCreationShooting home_defencePressure
## 1             69            56            36
## 2             57            47            51
## 3             53            47            53
## 4             42            53            42
## 5             50            48            39
## 6             50            50            45
##   home_defenceAggression home_defenceTeamWidth away_team_long_name
## 1             57            70            KV Oostende
## 2             50            57            KV Mechelen
## 3             50            61            Lierse SK
## 4             47            44            KV Kortrijk
## 5             48            49            KSV Cercle Brugge
## 6             45            50            KAA Gent
##   away_team_short_name away_team_fifa_api_id away_buildUpPlaySpeed
## 1           OOS          682          52
## 2           MEC         110724         52
## 3           LIE          239          69
## 4           KOR         100081         42
## 5           CEB          1750         53
## 6           GEN          674          50
##   away_buildUpPlayDribbling away_buildUpPlayPassing away_chanceCreationPassing
## 1             49            54            45
## 2             42            40            47
## 3             44            58            56
## 4             40            47            37
## 5             49            43            52
## 6             55            37            42
##   away_chanceCreationCrossing away_chanceCreationShooting away_defencePressure
## 1             42            53            42
## 2             53            50            42
## 3             52            61            38
## 4             67            30            46
## 5             58            42            59

```

```

## 6          66          34          49
##   away_defenceAggression away_defenceTeamWidth
## 1           47           44
## 2           51           52
## 3           49           65
## 4           45           52
## 5           47           62
## 6           54           52

#write.csv(train_matches, file = "train_matches.csv", row.names = FALSE)
#write.csv(test_matches, file = "test_matches.csv", row.names = FALSE)

# Prepare training and test data with predictors only
train_data <- train_matches |>
  mutate(
    buildUpPlaySpeed_diff      = home_buildUpPlaySpeed      - away_buildUpPlaySpeed,
    buildUpPlayPassing_diff     = home_buildUpPlayPassing     - away_buildUpPlayPassing,
    chanceCreationPassing_diff = home_chanceCreationPassing - away_chanceCreationPassing,
    chanceCreationCrossing_diff = home_chanceCreationCrossing - away_chanceCreationCrossing,
    chanceCreationShooting_diff = home_chanceCreationShooting - away_chanceCreationShooting,
    defencePressure_diff        = home_defencePressure        - away_defencePressure,
    defenceAggression_diff     = home_defenceAggression     - away_defenceAggression,
    defenceTeamWidth_diff       = home_defenceTeamWidth       - away_defenceTeamWidth,
    speed_shooting_int          = buildUpPlaySpeed_diff * chanceCreationShooting_diff,
    defence_press_agg_int       = defencePressure_diff * defenceAggression_diff
  ) |>
  select(
    home_result,
    buildUpPlaySpeed_diff,
    buildUpPlayPassing_diff,
    chanceCreationPassing_diff,
    chanceCreationCrossing_diff,
    chanceCreationShooting_diff,
    defencePressure_diff,
    defenceAggression_diff,
    defenceTeamWidth_diff,
    speed_shooting_int,
    defence_press_agg_int
  ) |>
  drop_na()

test_data <- test_matches |>
  mutate(
    buildUpPlaySpeed_diff      = home_buildUpPlaySpeed      - away_buildUpPlaySpeed,
    buildUpPlayPassing_diff     = home_buildUpPlayPassing     - away_buildUpPlayPassing,
    chanceCreationPassing_diff = home_chanceCreationPassing - away_chanceCreationPassing,
    chanceCreationCrossing_diff = home_chanceCreationCrossing - away_chanceCreationCrossing,
    chanceCreationShooting_diff = home_chanceCreationShooting - away_chanceCreationShooting,
    defencePressure_diff        = home_defencePressure        - away_defencePressure,
    defenceAggression_diff     = home_defenceAggression     - away_defenceAggression,
    defenceTeamWidth_diff       = home_defenceTeamWidth       - away_defenceTeamWidth,
    speed_shooting_int          = buildUpPlaySpeed_diff * chanceCreationShooting_diff,
    defence_press_agg_int       = defencePressure_diff * defenceAggression_diff
  ) |>
  select(

```

```

    home_result,
    buildUpPlaySpeed_diff,
    buildUpPlayPassing_diff,
    chanceCreationPassing_diff,
    chanceCreationCrossing_diff,
    chanceCreationShooting_diff,
    defencePressure_diff,
    defenceAggression_diff,
    defenceTeamWidth_diff,
    speed_shooting_int,
    defence_press_agg_int
  )

## Drop predictors that are entirely NA ----
na_all_cols <- names(train_data)[sapply(train_data, function(x) all(is.na(x)))]  

na_all_cols

## character(0)
train_data <- train_data |>
  select(-all_of(na_all_cols))

test_data <- test_data |>
  select(-all_of(na_all_cols))

## Also drop any remaining rows with NA in predictors or outcome
train_data <- train_data |>
  drop_na()

test_data <- test_data |>
  drop_na()

## Logistic Regression ----

# Specify multinomial logistic regression
logistic_spec <- multinom_reg(mode = "classification") |>
  set_engine("nnet")

# Fit model on training data
logistic_spec |>
  fit(home_result ~ ., data = train_data) -> m0.fit

# Training confusion matrix
cat("==> MULTINOMIAL LOGISTIC RESULTS ==>\n")

## ==> MULTINOMIAL LOGISTIC RESULTS ==
m0.fit |>
  augment(new_data = train_data) |>
  conf_mat(truth = home_result, estimate = .pred_class)

##           Truth
## Prediction Win Draw Loss
##       Win   6607 3539 3854
##       Draw      0     0     0

```

```

##      Loss  369  283  438
# Training error rate
m0.fit |>
  augment(new_data = train_data) |>
  accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.5330683

# Test confusion matrix and error rate
logistic_spec |>
  fit(home_result ~ ., data = train_data) |>
  augment(new_data = test_data) -> m0.test_res

m0.test_res |>
  conf_mat(truth = home_result, estimate = .pred_class)

##          Truth
## Prediction Win Draw Loss
##      Win    1337   743   850
##      Draw      0     0     0
##      Loss     68    63    92

m0.test_res |>
  accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.5467808

#### LDA ----

cat("== LDA RESULTS ==\n")

## == LDA RESULTS ==
# Specify LDA model
lda_spec <- discrim_linear()

# Fit on training data
lda_spec |>
  fit(home_result ~ ., data = train_data) -> lda.fit

# Training confusion matrix
lda.fit |>
  augment(new_data = train_data) |>
  conf_mat(truth = home_result, estimate = .pred_class)

##          Truth
## Prediction Win Draw Loss
##      Win    6613  3544  3858
##      Draw      0     0     0
##      Loss     363   278   434

# Training error rate
lda.fit |>
  augment(new_data = train_data) |>

```

```

accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.533002

# Test confusion matrix + test error
lda.fit |>
  augment(new_data = test_data) -> lda.test_res

lda.test_res |>
  conf_mat(truth = home_result, estimate = .pred_class)

##          Truth
## Prediction Win Draw Loss
##      Win    1337   745   847
##      Draw      0     0     0
##      Loss     68    61    95

lda.test_res |>
  accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.5458294

### QDA ----

cat("==== QDA RESULTS ====\n")

## === QDA RESULTS ===

qda_spec <- discrim_quad()

# Fit on training data
qda_spec |>
  fit(home_result ~ ., data = train_data) -> qda.fit

# Training confusion matrix
qda.fit |>
  augment(new_data = train_data) |>
  conf_mat(truth = home_result, estimate = .pred_class)

##          Truth
## Prediction Win Draw Loss
##      Win    5986  3134  3343
##      Draw     35    31    30
##      Loss    955   657   919

# Training error
qda.fit |>
  augment(new_data = train_data) |>
  accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.5404904

```

```

# Test predictions
qda.fit |>
  augment(new_data = test_data) -> qda.test_res

qda.test_res |>
  conf_mat(truth = home_result, estimate = .pred_class)

##          Truth
## Prediction Win Draw Loss
##      Win    1214   648   704
##      Draw     5     3     0
##      Loss    186   155   238

qda.test_res |>
  accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.5385347

#### KNN ----

cat("==== KNN (K=1) RESULTS ====\n")

## === KNN (K=1) RESULTS ===

## K = 1
knn1_spec <- nearest_neighbor(mode = "classification", neighbors = 1)

knn1_spec |>
  fit(home_result ~ ., data = train_data) -> knn1.fit

knn1.fit |>
  augment(new_data = test_data) -> knn1.test_res

knn1.test_res |>
  conf_mat(truth = home_result, estimate = .pred_class)

##          Truth
## Prediction Win Draw Loss
##      Win    681   361   421
##      Draw    356   201   241
##      Loss    368   244   280

knn1.test_res |>
  accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.6314621

## K = 3

cat("==== KNN (K=3) RESULTS ====\n")

## === KNN (K=3) RESULTS ===

```

```

knn3_spec <- nearest_neighbor(mode = "classification", neighbors = 3)

knn3_spec |>
  fit(home_result ~ ., data = train_data) -> knn3.fit

knn3.fit |>
  augment(new_data = test_data) -> knn3.test_res

knn3.test_res |>
  conf_mat(truth = home_result, estimate = .pred_class)

##          Truth
## Prediction Win Draw Loss
##      Win   678   359   420
##      Draw   364   199   245
##      Loss   363   248   277

knn3.test_res |>
  accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.6339994
## K = 5

cat("== KNN (K=5) RESULTS ==\n")

## == KNN (K=5) RESULTS ==

knn5_spec <- nearest_neighbor(mode = "classification", neighbors = 5)

knn5_spec |>
  fit(home_result ~ ., data = train_data) -> knn5.fit

knn5.fit |>
  augment(new_data = test_data) -> knn5.test_res

knn5.test_res |>
  conf_mat(truth = home_result, estimate = .pred_class)

##          Truth
## Prediction Win Draw Loss
##      Win   738   390   429
##      Draw   327   189   227
##      Loss   340   227   286

knn5.test_res |>
  accuracy(truth = home_result, estimate = .pred_class) |>
  mutate(error = 1 - .estimate) |>
  pull(error)

## [1] 0.615287

### Validation Set Approach ----

goals_train <- train_matches |>

```

```

mutate(
  buildUpPlaySpeed_diff      = home_buildUpPlaySpeed           - away_buildUpPlaySpeed,
  buildUpPlayPassing_diff     = home_buildUpPlayPassing        - away_buildUpPlayPassing,
  chanceCreationPassing_diff = home_chanceCreationPassing    - away_chanceCreationPassing,
  chanceCreationCrossing_diff = home_chanceCreationCrossing   - away_chanceCreationCrossing,
  chanceCreationShooting_diff = home_chanceCreationShooting  - away_chanceCreationShooting,
  defencePressure_diff        = home_defencePressure          - away_defencePressure,
  defenceAggression_diff     = home_defenceAggression        - away_defenceAggression,
  defenceTeamWidth_diff       = home_defenceTeamWidth         - away_defenceTeamWidth,
  speed_shooting_int          = buildUpPlaySpeed_diff * chanceCreationShooting_diff,
  defence_press_agg_int       = defencePressure_diff * defenceAggression_diff
) |>
select(
  home_team_goal,
  buildUpPlaySpeed_diff,
  buildUpPlayPassing_diff,
  chanceCreationPassing_diff,
  chanceCreationCrossing_diff,
  chanceCreationShooting_diff,
  defencePressure_diff,
  defenceAggression_diff,
  defenceTeamWidth_diff,
  speed_shooting_int,
  defence_press_agg_int
)

goals_test <- test_matches |>
  mutate(
    buildUpPlaySpeed_diff      = home_buildUpPlaySpeed           - away_buildUpPlaySpeed,
    buildUpPlayPassing_diff     = home_buildUpPlayPassing        - away_buildUpPlayPassing,
    chanceCreationPassing_diff = home_chanceCreationPassing    - away_chanceCreationPassing,
    chanceCreationCrossing_diff = home_chanceCreationCrossing   - away_chanceCreationCrossing,
    chanceCreationShooting_diff = home_chanceCreationShooting  - away_chanceCreationShooting,
    defencePressure_diff        = home_defencePressure          - away_defencePressure,
    defenceAggression_diff     = home_defenceAggression        - away_defenceAggression,
    defenceTeamWidth_diff       = home_defenceTeamWidth         - away_defenceTeamWidth,
    speed_shooting_int          = buildUpPlaySpeed_diff * chanceCreationShooting_diff,
    defence_press_agg_int       = defencePressure_diff * defenceAggression_diff
) |>
select(
  home_team_goal,
  buildUpPlaySpeed_diff,
  buildUpPlayPassing_diff,
  chanceCreationPassing_diff,
  chanceCreationCrossing_diff,
  chanceCreationShooting_diff,
  defencePressure_diff,
  defenceAggression_diff,
  defenceTeamWidth_diff,
  speed_shooting_int,
  defence_press_agg_int
)

```

```

## Drop predictors that are entirely NA ----
na_all_cols2 <- names(goals_train) [sapply(goals_train, function(x) all(is.na(x)))] 

goals_train <- goals_train |>
  select(-all_of(na_all_cols2))

goals_test <- goals_test |>
  select(-all_of(na_all_cols2))

## Also drop any remaining rows with NA in predictors or outcome
goals_train <- goals_train |>
  drop_na()

goals_test <- goals_test |>
  drop_na()

## Prepare linear regression spec

linear_spec <- linear_reg()

## Base recipe: predict home_team_goal using the best predictors
linear_rec <- recipe(
  home_team_goal ~ ., data = goals_train)

## 4. Fit the linear model
linear_model <- workflow() |>
  add_model(linear_spec) |>
  add_recipe(linear_rec) |>
  fit(data = goals_train)

## Compute test MSE for linear model

linear_model |>
  augment(new_data = goals_test) |>
  mutate(resid2 = (.resid)^2) |>
  summarise(mse = mean(resid2, na.rm = TRUE)) |>
  mutate(model = "Linear Regression (Goals)") |>
  relocate(model)

## # A tibble: 1 x 2
##   model           mse
##   <chr>        <dbl>
## 1 Linear Regression (Goals) 1.64
### Quadratic Model ----

## Add squared terms

quad_rec <- linear_rec |>
  step_mutate(
    buildUpPlaySpeed_diff2      = buildUpPlaySpeed_diff^2,
    buildUpPlayPassing_diff2    = buildUpPlayPassing_diff^2,
    chanceCreationPassing_diff2 = chanceCreationPassing_diff^2,

```

```

chanceCreationCrossing_diff2 = chanceCreationCrossing_diff^2,
chanceCreationShooting_diff2 = chanceCreationShooting_diff^2,
defencePressure_diff2       = defencePressure_diff^2,
defenceAggression_diff2    = defenceAggression_diff^2,
defenceTeamWidth_diff2      = defenceTeamWidth_diff^2
)

workflow() |>
  add_model(linear_spec) |>
  add_recipe(quad_rec) |>
  fit(data = goals_train) |>
  augment(new_data = goals_test) |>
  mutate(resid2 = (.resid)^2) |>
  summarise(mse = mean(resid2, na.rm = TRUE)) |>
  mutate(model = "Quadratic Regression (Goals)") |>
  relocate(model)

## # A tibble: 1 x 2
##   model                 mse
##   <chr>                <dbl>
## 1 Quadratic Regression (Goals) 1.64

### Cubic Model ----

## Add cubic terms the same way

cubic_rec <- quad_rec |>
  step_mutate(
    buildUpPlaySpeed_diff3      = buildUpPlaySpeed_diff^3,
    buildUpPlayPassing_diff3    = buildUpPlayPassing_diff^3,
    chanceCreationPassing_diff3 = chanceCreationPassing_diff^3,
    chanceCreationCrossing_diff3 = chanceCreationCrossing_diff^3,
    chanceCreationShooting_diff3 = chanceCreationShooting_diff^3,
    defencePressure_diff3       = defencePressure_diff^3,
    defenceAggression_diff3    = defenceAggression_diff^3,
    defenceTeamWidth_diff3      = defenceTeamWidth_diff^3
  )

workflow() |>
  add_model(linear_spec) |>
  add_recipe(cubic_rec) |>
  fit(data = goals_train) |>
  augment(new_data = goals_test) |>
  mutate(resid2 = (.resid)^2) |>
  summarise(mse = mean(resid2, na.rm = TRUE)) |>
  mutate(model = "Cubic Regression (Goals)") |>
  relocate(model)

## # A tibble: 1 x 2
##   model                 mse
##   <chr>                <dbl>
## 1 Cubic Regression (Goals) 1.63

### KNN ----

```

```

goals_knn <- train_matches |>
  mutate(
    buildUpPlaySpeed_diff      = home_buildUpPlaySpeed      - away_buildUpPlaySpeed,
    buildUpPlayPassing_diff     = home_buildUpPlayPassing     - away_buildUpPlayPassing,
    chanceCreationPassing_diff = home_chanceCreationPassing - away_chanceCreationPassing,
    chanceCreationCrossing_diff = home_chanceCreationCrossing - away_chanceCreationCrossing,
    chanceCreationShooting_diff = home_chanceCreationShooting - away_chanceCreationShooting,
    defencePressure_diff       = home_defencePressure       - away_defencePressure,
    defenceAggression_diff     = home_defenceAggression     - away_defenceAggression,
    defenceTeamWidth_diff      = home_defenceTeamWidth      - away_defenceTeamWidth,
    speed_shooting_int          = buildUpPlaySpeed_diff * chanceCreationShooting_diff,
    defence_press_agg_int       = defencePressure_diff * defenceAggression_diff
  ) |>
  select(
    home_team_goal,
    buildUpPlaySpeed_diff,
    buildUpPlayPassing_diff,
    chanceCreationPassing_diff,
    chanceCreationCrossing_diff,
    chanceCreationShooting_diff,
    defencePressure_diff,
    defenceAggression_diff,
    defenceTeamWidth_diff,
    speed_shooting_int,
    defence_press_agg_int
  ) |>
  drop_na()

na_all_cols3 <- names(goals_knn)[sapply(goals_knn, function(x) all(is.na(x)))] 

goals_knn <- goals_knn |>
  select(-all_of(na_all_cols3))

## Also drop any remaining rows with NA in predictors or outcome
goals_knn <- goals_knn |>
  drop_na()

## KNN k-fold CV function ----
k_fold_cv_err_knn <- function(k_fold = 10, knn, data, formula) {
  data.kfold <- vfold_cv(data, v = k_fold)

  knn_spec <- nearest_neighbor(mode = "regression", neighbors = knn)

  knn_rec <- recipe(formula, data = data) |>
    step_normalize(all_predictors())

  workflow() |>
    add_model(knn_spec) |>
    add_recipe(knn_rec) |>
    fit_resamples(data.kfold) |>
    collect_metrics() |>
    select(.metric, mean) |>
    pivot_wider(names_from = .metric, values_from = mean) |>

```

```

    mutate(mse = rmse^2) |>
    pull(mse)
}

## Evaluate KNN for K = 1, 5, 10, 20, 100 ----
res <- data.frame(knn = c(1, 5, 10, 20, 100))

for(i in seq_len(nrow(res))) {
  res[i, "mse"] <- k_fold_cv_err_knn(
    k_fold = 10,
    knn = res[i, "knn"],
    data = goals_knn,
    formula = home_team_goal ~ .
  )
}

## Inspect results and best K ----
res

##      knn      mse
## 1     1 2.986017
## 2     5 2.038056
## 3    10 1.826022
## 4    20 1.729794
## 5   100 1.663585
res[which.min(res$mse), ]

##      knn      mse
## 5 100 1.663585

# Build complete data set with engineered predictors and no missing values

goals_ridge <- train_matches |>
  mutate(
    buildUpPlaySpeed_diff      = home_buildUpPlaySpeed      - away_buildUpPlaySpeed,
    buildUpPlayPassing_diff     = home_buildUpPlayPassing     - away_buildUpPlayPassing,
    chanceCreationPassing_diff = home_chanceCreationPassing - away_chanceCreationPassing,
    chanceCreationCrossing_diff = home_chanceCreationCrossing - away_chanceCreationCrossing,
    chanceCreationShooting_diff = home_chanceCreationShooting - away_chanceCreationShooting,
    defencePressure_diff        = home_defencePressure        - away_defencePressure,
    defenceAggression_diff     = home_defenceAggression     - away_defenceAggression,
    defenceTeamWidth_diff       = home_defenceTeamWidth       - away_defenceTeamWidth,
    speed_shooting_int          = buildUpPlaySpeed_diff * chanceCreationShooting_diff,
    defence_press_agg_int       = defencePressure_diff * defenceAggression_diff
  ) |>
  select(
    home_team_goal,
    buildUpPlaySpeed_diff,
    buildUpPlayPassing_diff,
    chanceCreationPassing_diff,
    chanceCreationCrossing_diff,
    chanceCreationShooting_diff,
    defencePressure_diff,
    defenceAggression_diff,

```

```

    defenceTeamWidth_diff,
    speed_shooting_int,
    defence_press_agg_int
) |>
drop_na()

## Ridge Regression ----
# Lambda grid
lambda <- 10^seq(-2, 10, length.out = 100)
tune_df <- data.frame(lambda = lambda)

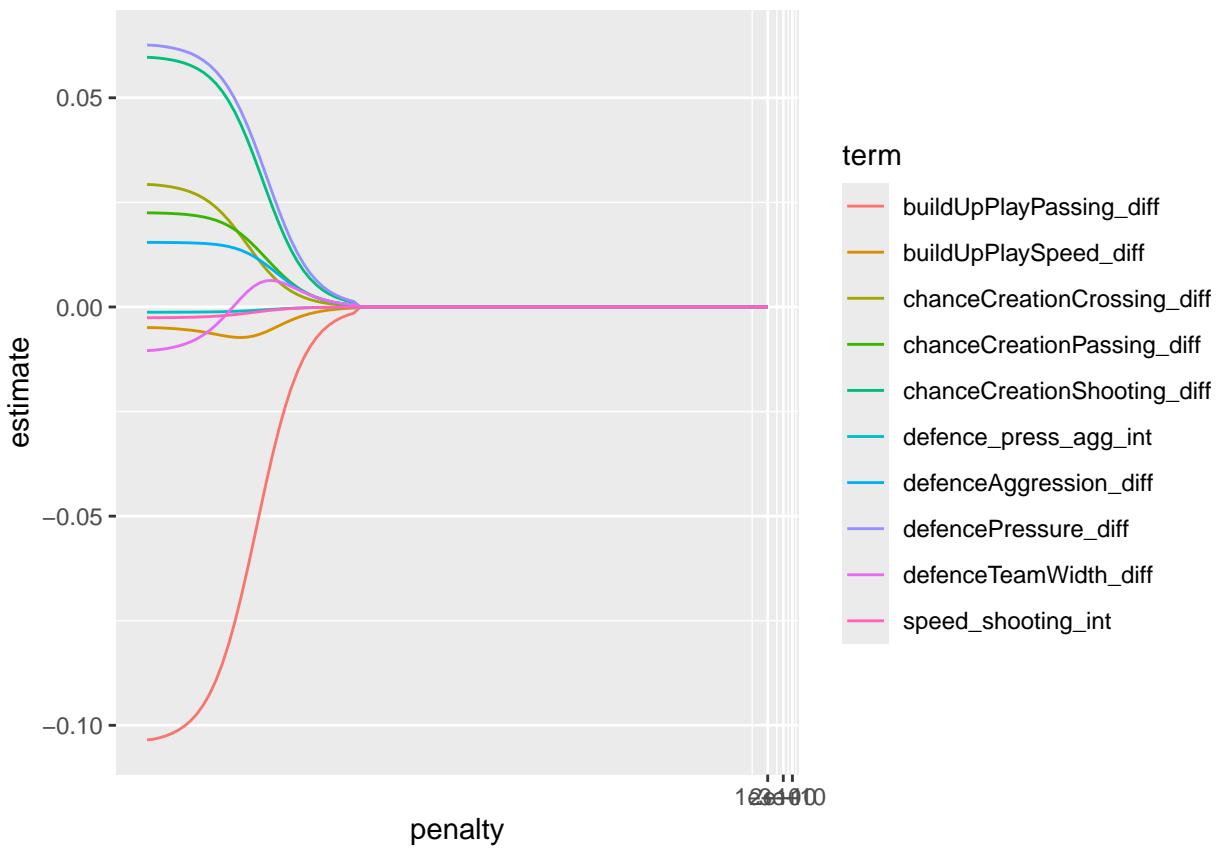
# Recipe: standardize predictors
prep_data <- recipe(home_team_goal ~ ., data = goals_ridge) |>
  step_normalize(all_predictors())

# Fit ridge for each lambda and store coefficient paths
ridge_ests <- data.frame()
for(lam in lambda) {
  ridge_spec <- linear_reg(mixture = 0, penalty = lam) |>
    set_mode("regression") |>
    set_engine("glmnet")

  workflow() |>
    add_model(ridge_spec) |>
    add_recipe(prep_data) |>
    fit(goals_ridge) |>
    tidy() |>
    bind_rows(ridge_ests) -> ridge_ests
}

##
## Attaching package: 'Matrix'
## The following objects are masked from 'package:tidyর':
##       expand, pack, unpack
## Loaded glmnet 4.1-10
# Plot coefficient paths vs lambda
ridge_ests |>
  filter(term != "(Intercept)") |>
  ggplot() +
  geom_line(aes(penalty, estimate, group = term, colour = term)) +
  coord_transform(x = "log10")

```

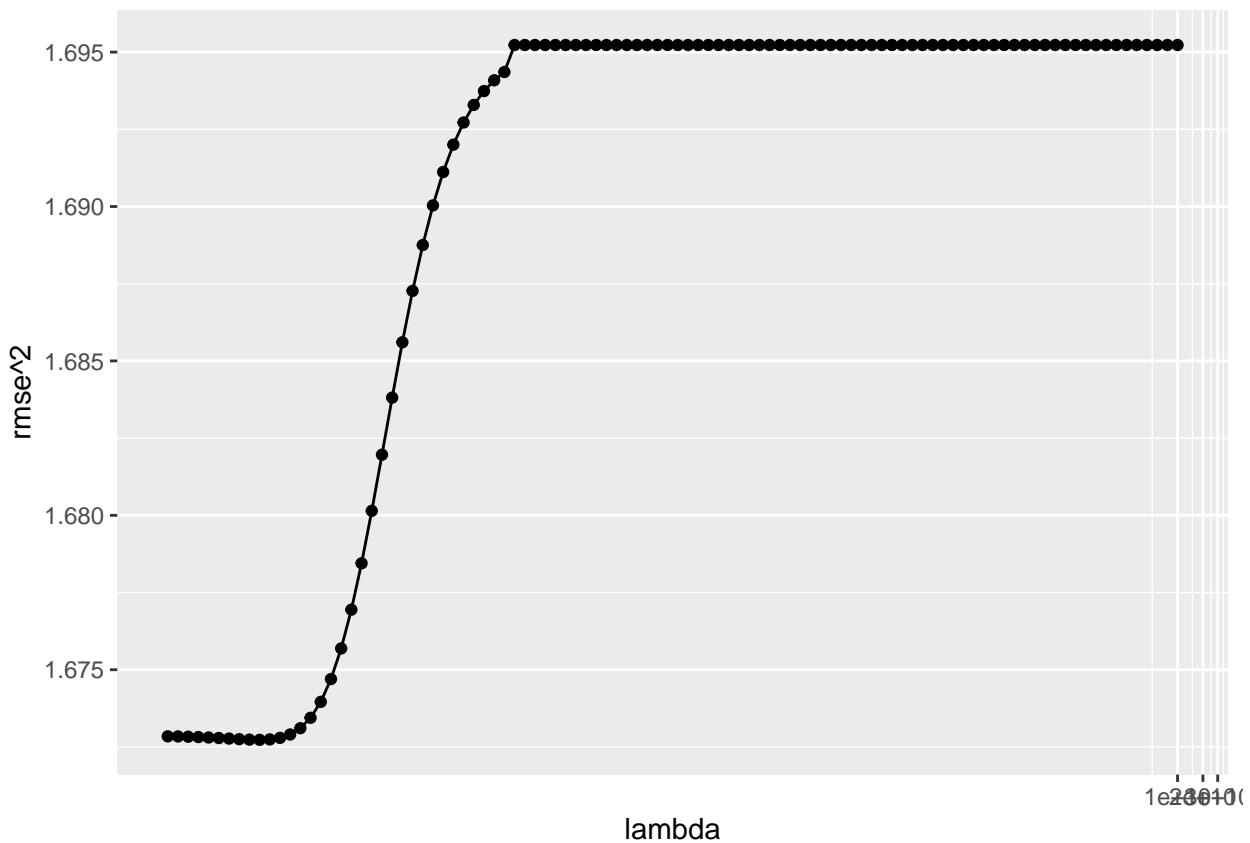


```
## 10-fold CV to choose lambda ----
goals_ridge_10foldcv <- vfold_cv(goals_ridge, v = 10)

ridge_spec <- linear_reg(mixture = 0, penalty = tune("lambda")) |>
  set_mode("regression") |>
  set_engine("glmnet")

workflow() |>
  add_model(ridge_spec) |>
  add_recipe(prep_data) |>
  tune_grid(resamples = goals_ridge_10foldcv, grid = tune_df, metrics = metric_set(rmse)) -> ridge_tune

ridge_tune |>
  collect_metrics() |>
  select(lambda, .metric, mean) |>
  pivot_wider(names_from = .metric, values_from = mean) |>
  ggplot() +
  geom_line(aes(lambda, rmse^2)) +
  geom_point(aes(lambda, rmse^2)) +
  coord_transform(x = "log10")
```



```
## Best penalty (lambda) based on CV RMSE
show_best(ridge_tune, metric = "rmse", n = 1)

## # A tibble: 1 x 7
##   lambda .metric .estimator  mean     n std_err .config
##   <dbl> <chr>    <chr>     <dbl> <int>  <dbl> <chr>
## 1  0.123 rmse     standard    1.29     10  0.0128 Preprocessor1_Model010
```