Predicting AirBnB Rental Prices

Group 11: Trevor Isaacson, Jonathan Olavarria, Jasmine DeMeyer

12/10/2021

Motivation

- You are looking for some additional income and decide renting on AirBnB is the best option
- How much should you rent your extra space for?

Data

- In general, AirBnB data is very open and be easily accessed
- The original dataset is from a past Kaggle competition
  - Contained over 74,000 individual listings
- For sake of time and processing power, we took a random sample of 17,500 from those 74,000 listings
- They also provided a testing file
- Since the competition is over, we will compile our final predictions on that file using our best model

Data

- Original data consists of 30 variables
- Variables are about the property, property location, the host and host reviews
- After cleaning and eliminating variables, our data consisted of 22 variables
- Property:
  - property_type, room_type, accommodates, bedrooms, beds, bed_type, bathrooms
- Location:
  - latitude, longitude, city
- Host:
  - cancellation_policy, cleaning_fee, host_has_profile_pic, host_identify_verified, etc

Baseline Regression

```
linear = lm(price ~ ., data = training)
```

```
## [1] "MSE of Testing Set:  0.165"
```

Regression Splines/Generalized Additive Models

PCR and PLS

- 10 Fold Cross-Validation was performed for number of components ranging from 1 to 20.
- The Cross-Validation MSE was used to pick optimal number of components for both models.

PCR

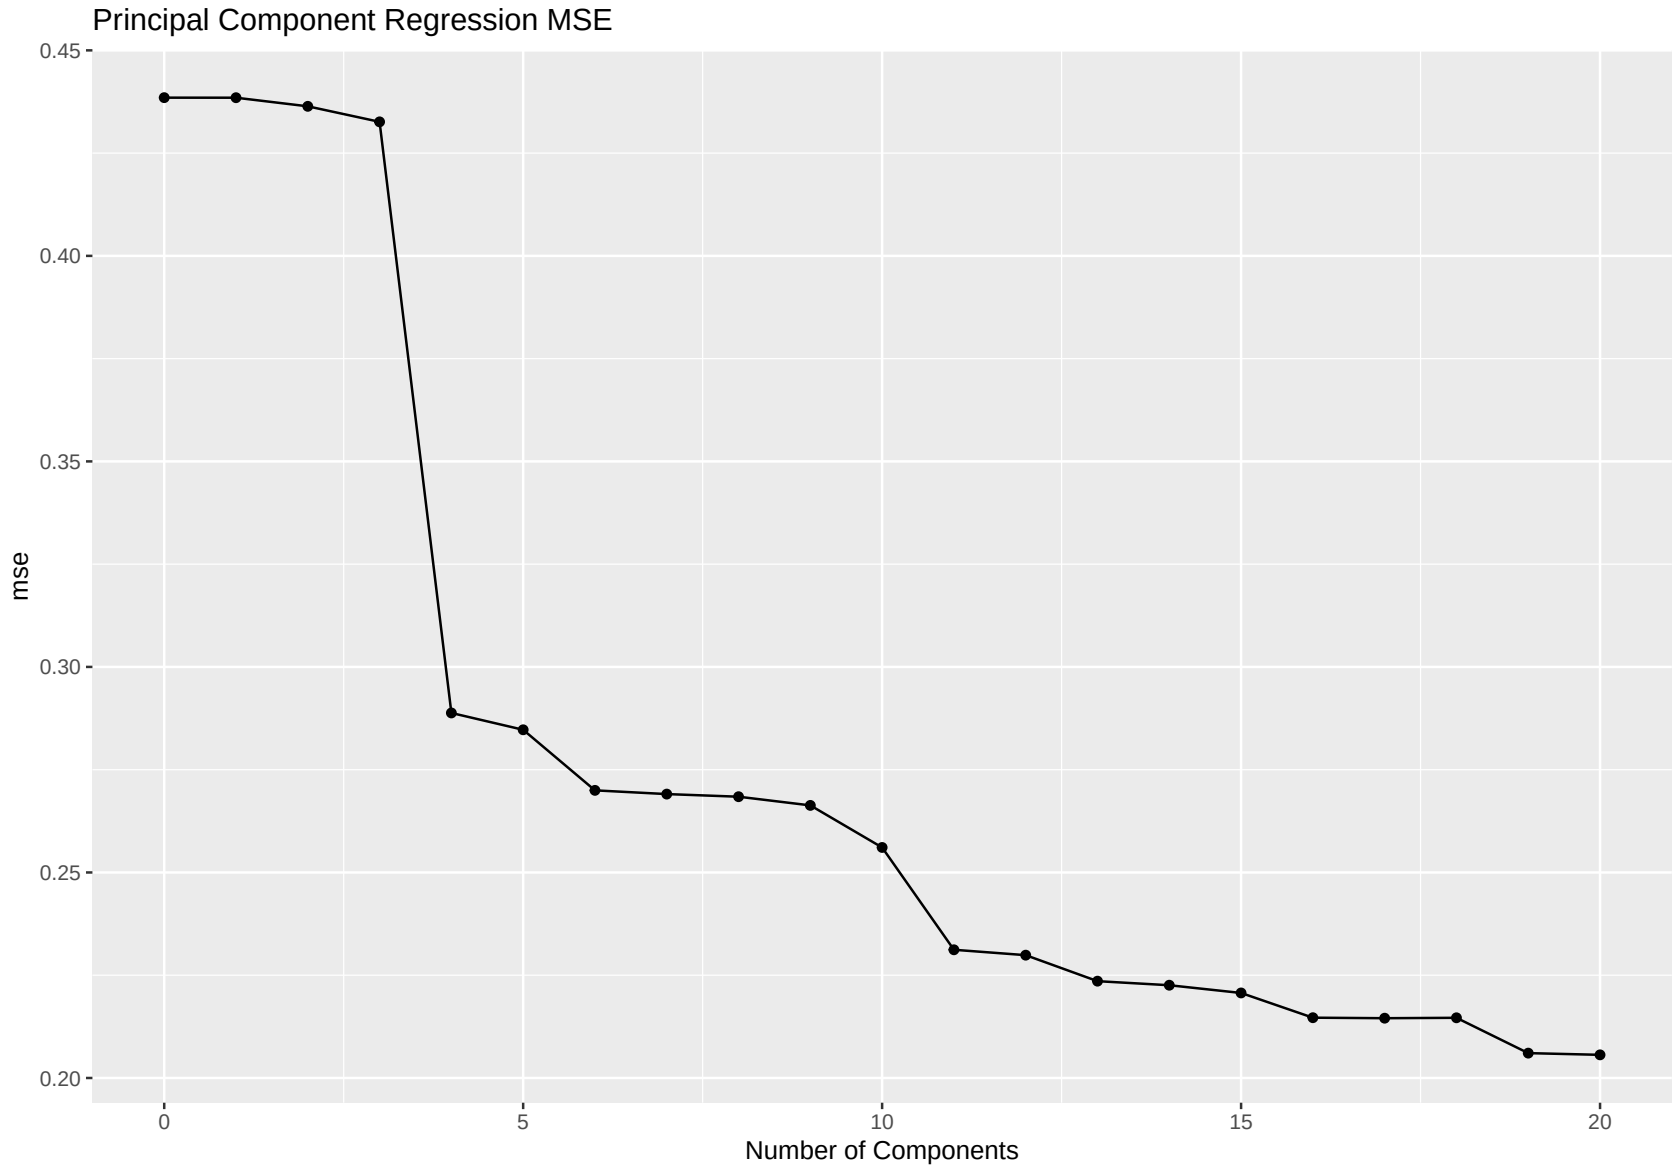
Principal Component Regression MSE
mse
0.45
0.40
0.35
0.30
0.25
0.20
0
5
10
15
20
Number of Components

PCR Predictions

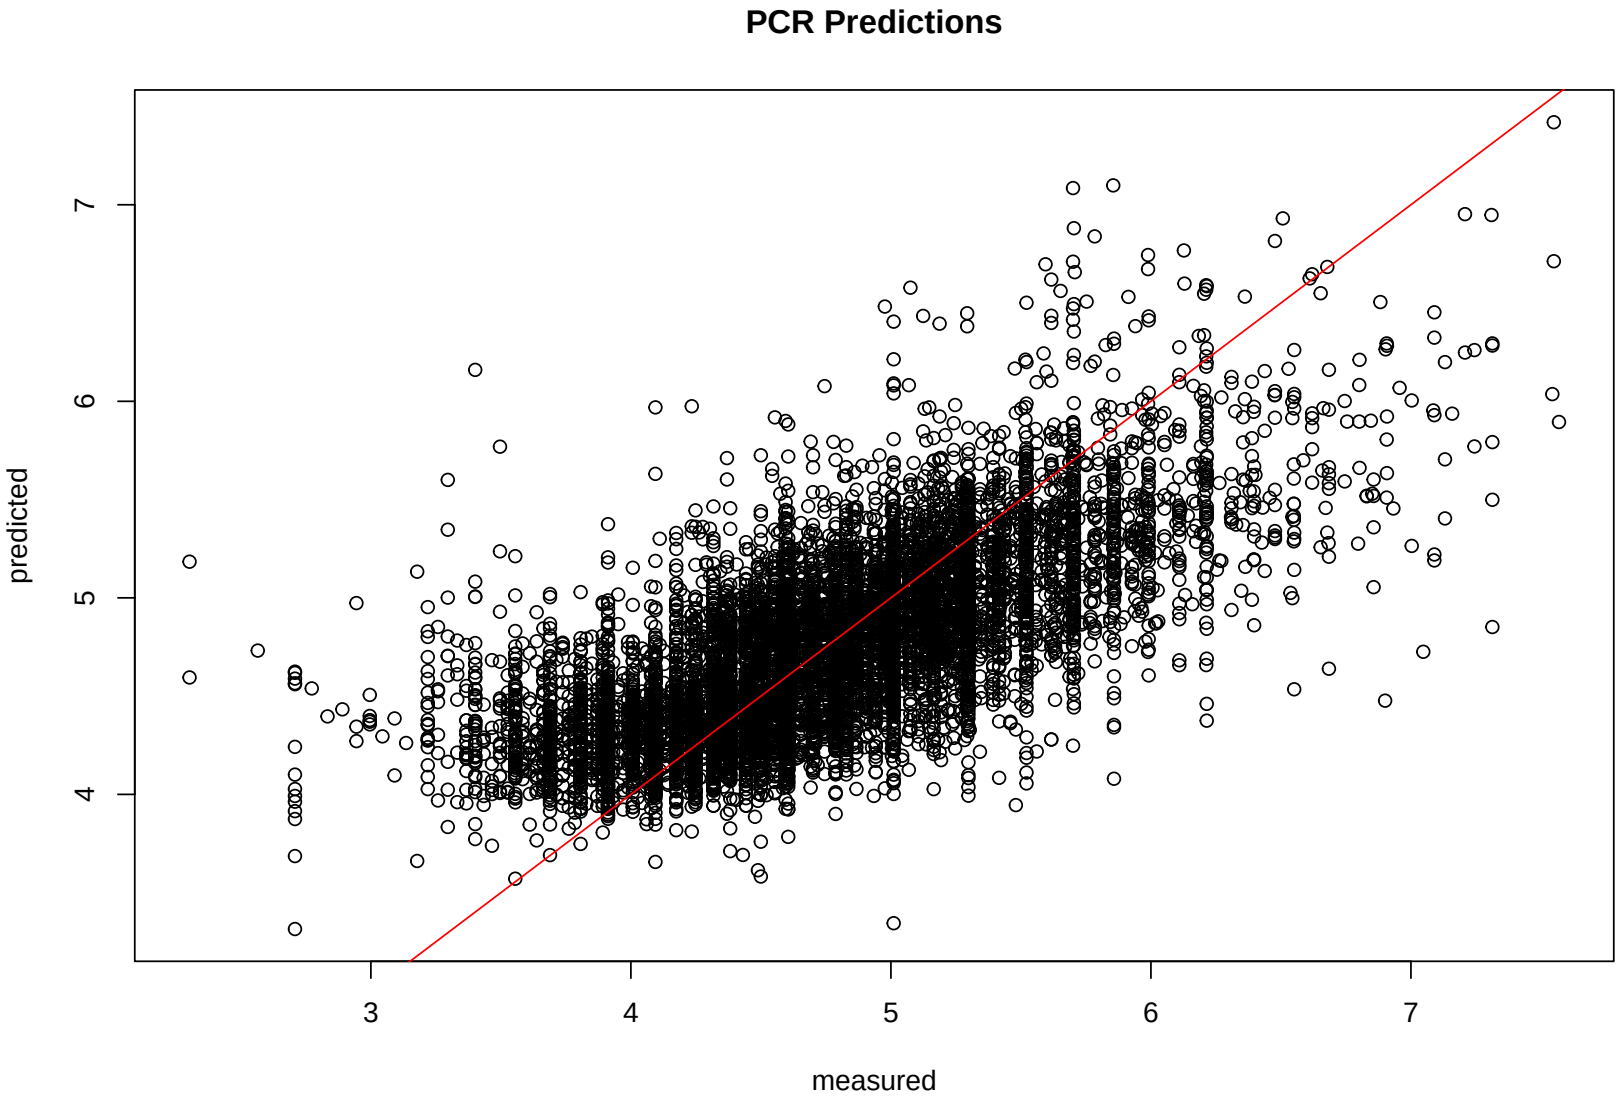

PLS

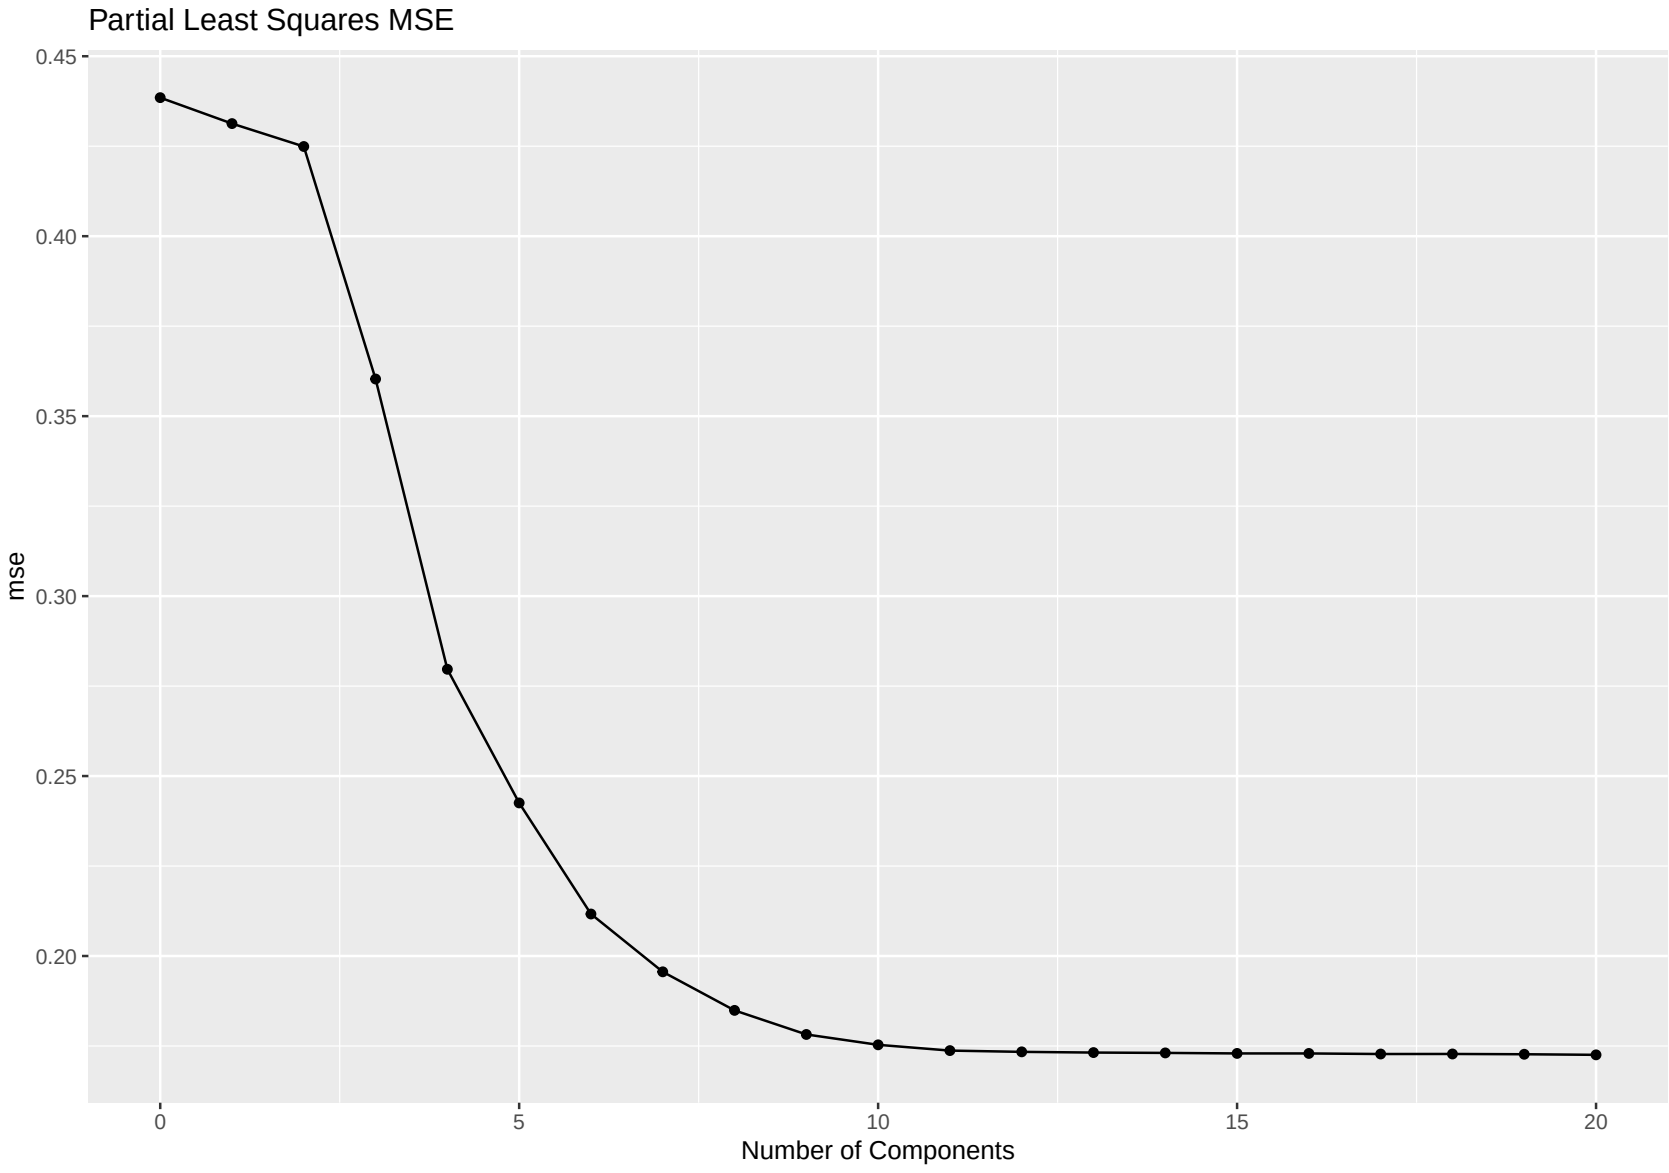
Partial Least Squares MSE
mse
0.45
0.40
0.35
0.30
0.25
0.20
0
5
10
15
20
Number of Components

PLS Predictions

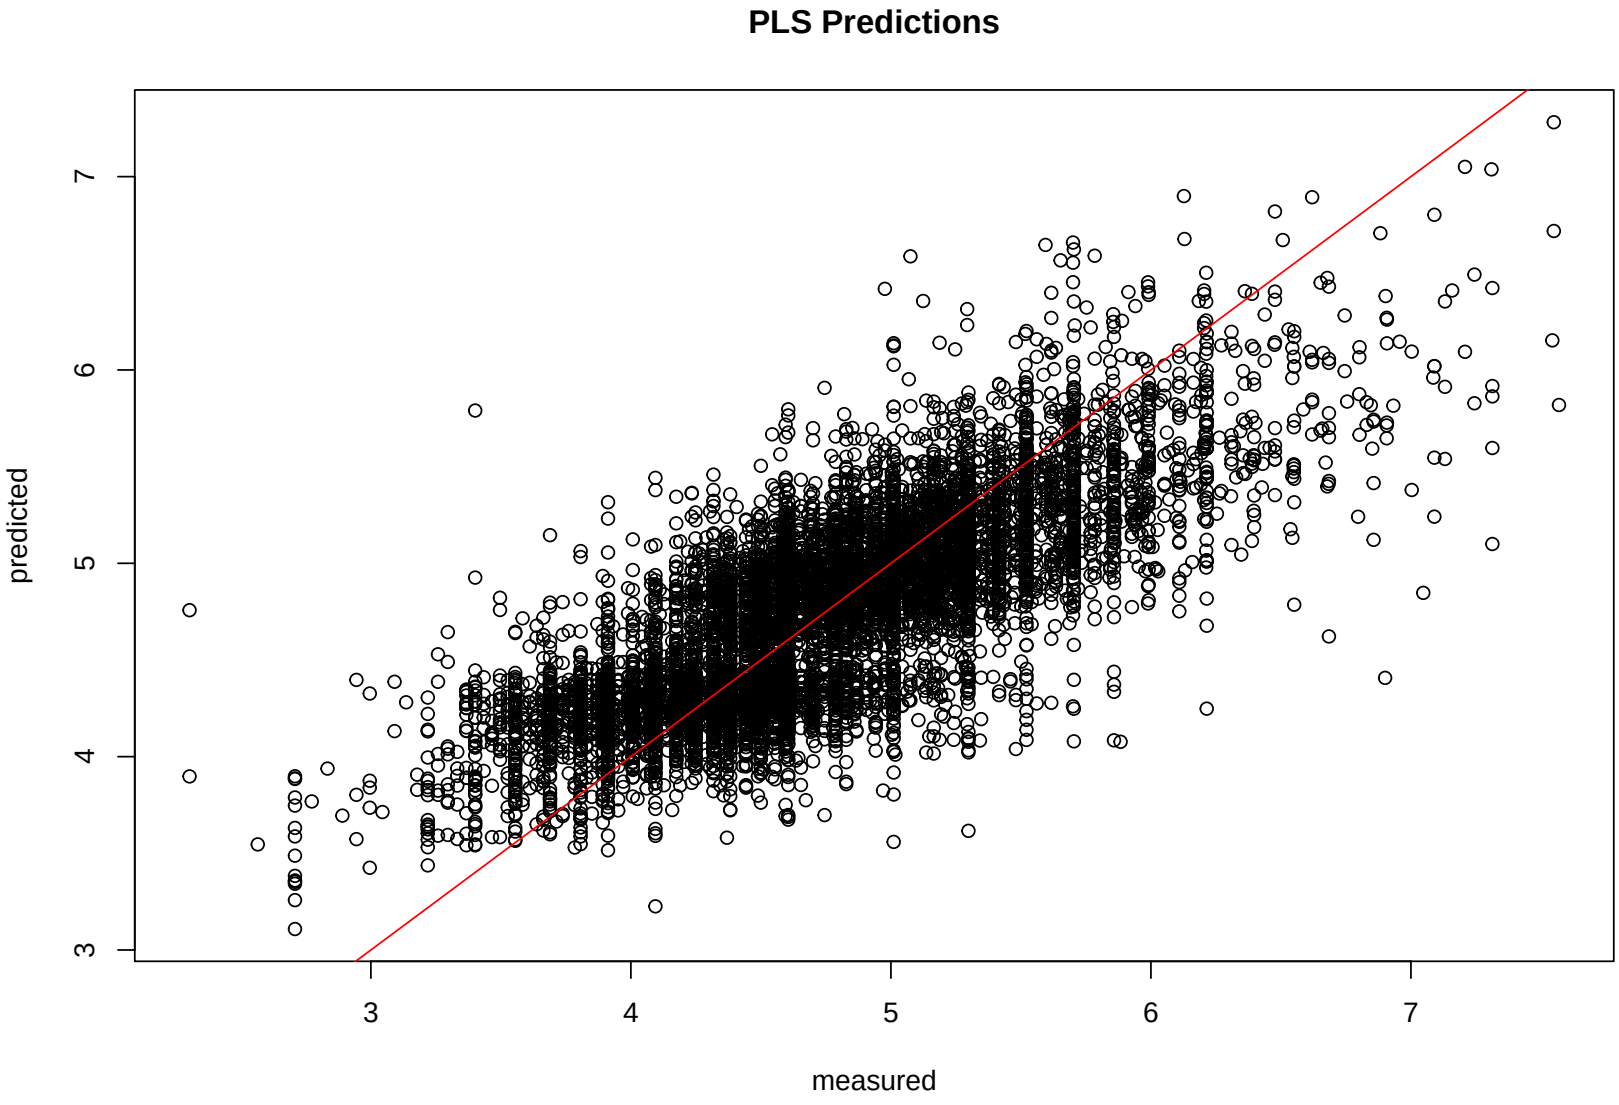

PCR and PLS Summary

```
##                              PCR     PLS
## Components                 15.0000 10.0000
## Test MSE                    0.1765  0.2192
## % Variance Explained       99.7000 99.9000
```

Regression Trees

```
##
## Regression tree:
## tree(formula = price ~ ., data = training)
## Variables actually used in tree construction:
## [1] "room_type" "longitude" "bathrooms" "city"        "bed
## Number of terminal nodes:  8
## Residual mean deviance:  0.1885 = 1695 / 8992
## Distribution of residuals:
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -2.5050 -0.2999 -0.0196  0.0000  0.2558  2.8310

## [1] "Test MSE of Initial Tree:  0.1926"
```

Regression Trees

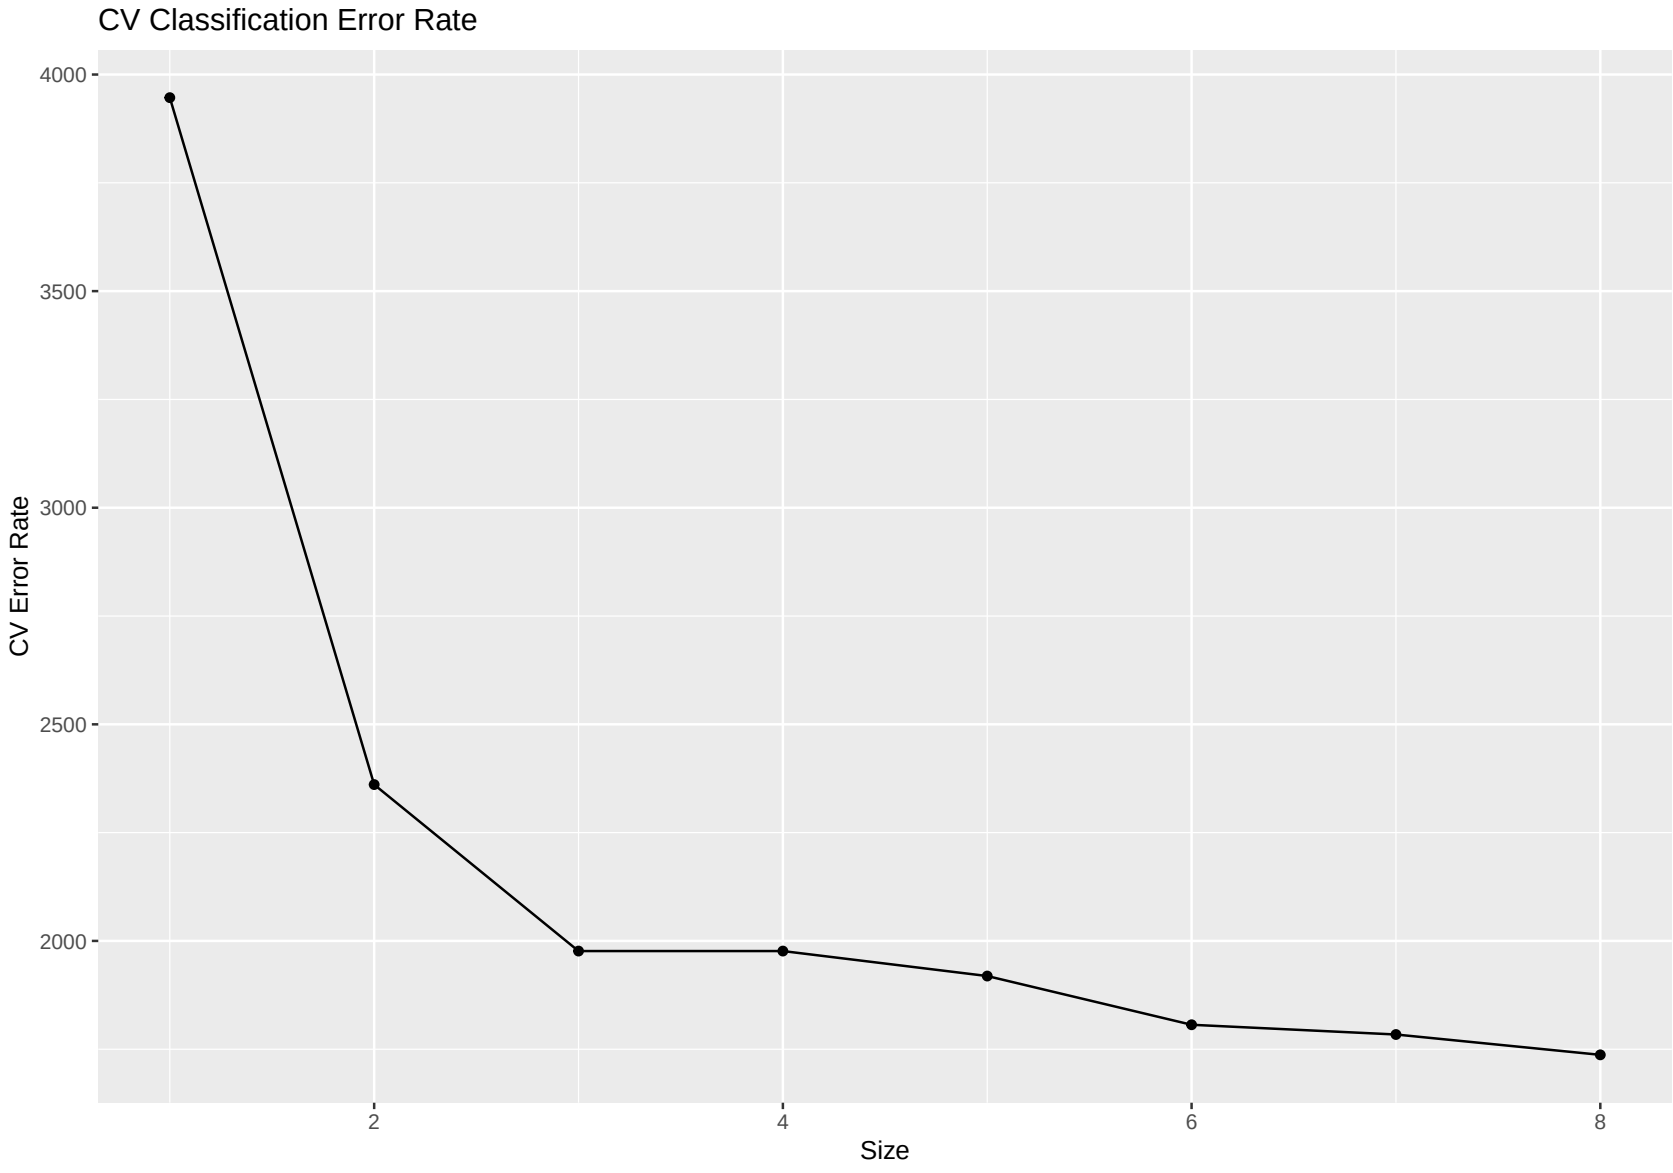

Regression Trees

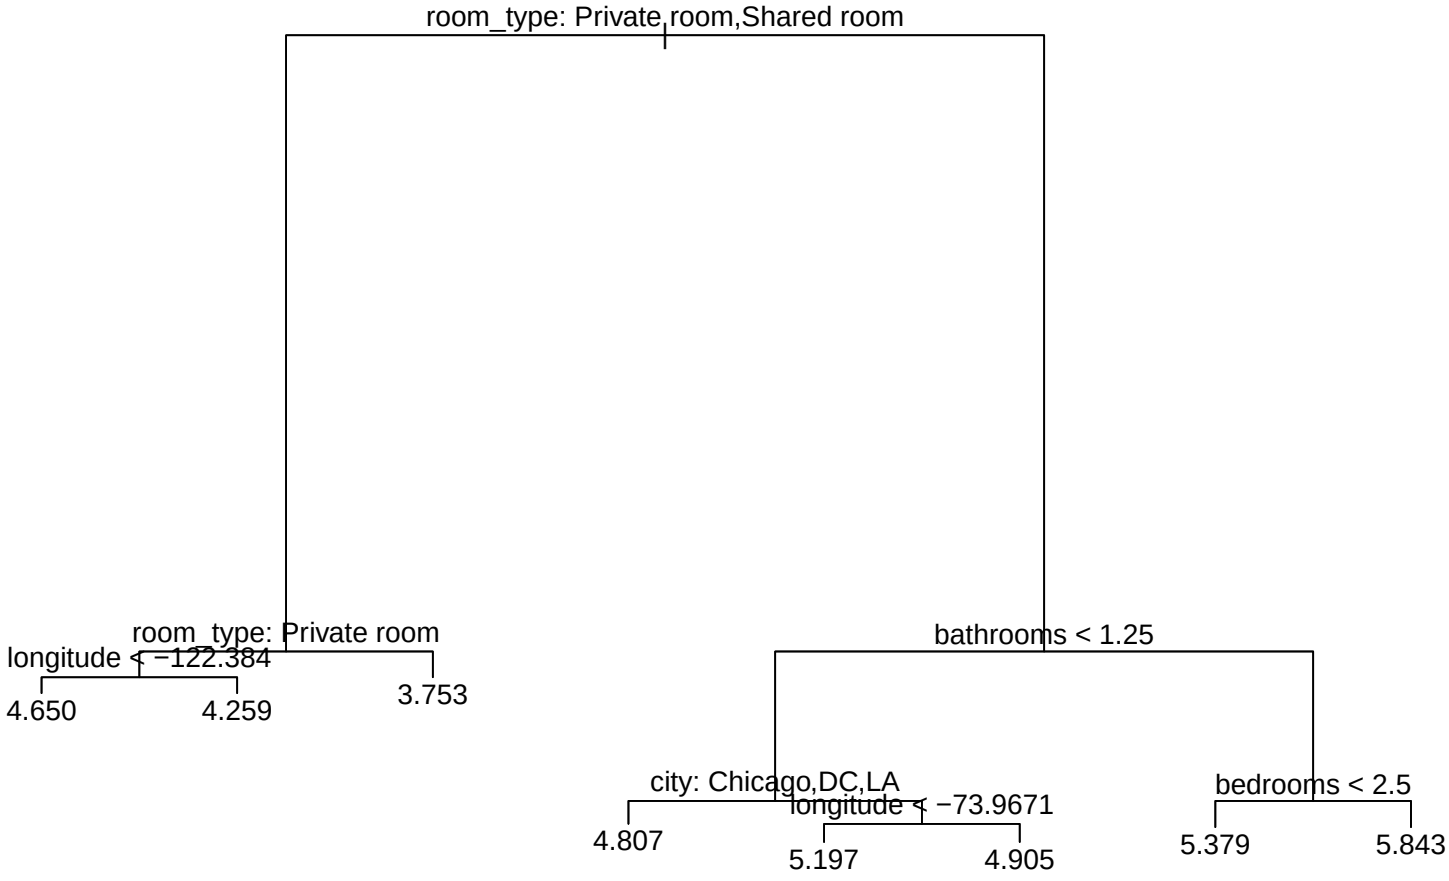

Bagging

```
bag_fit <- randomForest(price ~ ., data = training, mtry =
bag_predict = predict(bag_fit, testing, type = "response")
bag_MSE = round(mean((testing$price - bag_predict)^2), 4)
print(paste("Test MSE of Bagging: ", bag_MSE))
```

```
## [1] "Test MSE of Bagging:  0.1294"
```

Bagging

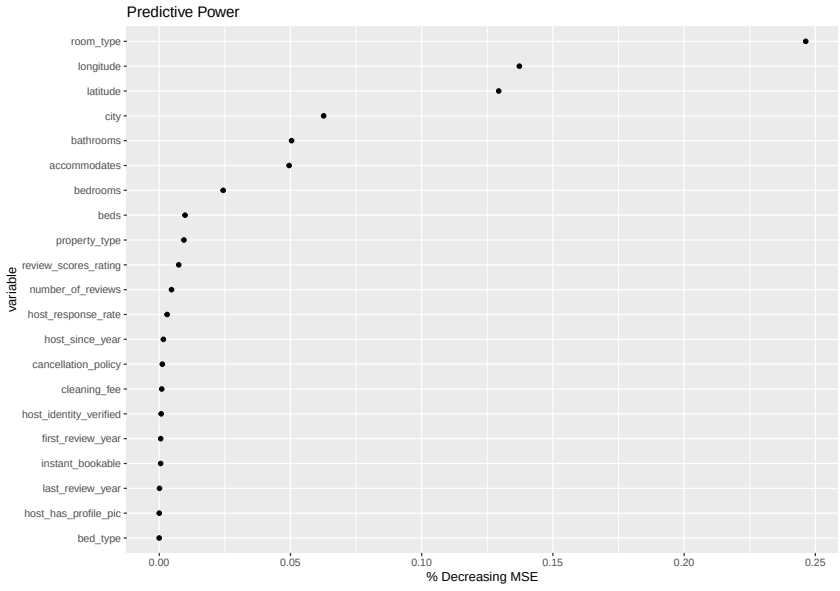
Predictive Power
room_type
longitude
latitude
city
bathrooms
accommodates
bedrooms
beds
property_type
review_scores_rating
number_of_reviews
host_response_rate
host_since_year
cancellation_policy
cleaning_fee
host_identity_verified
first_review_year
instant_bookable
last_review_year
host_has_profile_pic
bed_type
variable
0.00
0.05
0.10
0.15
0.20
0.25
% Decreasing MSE

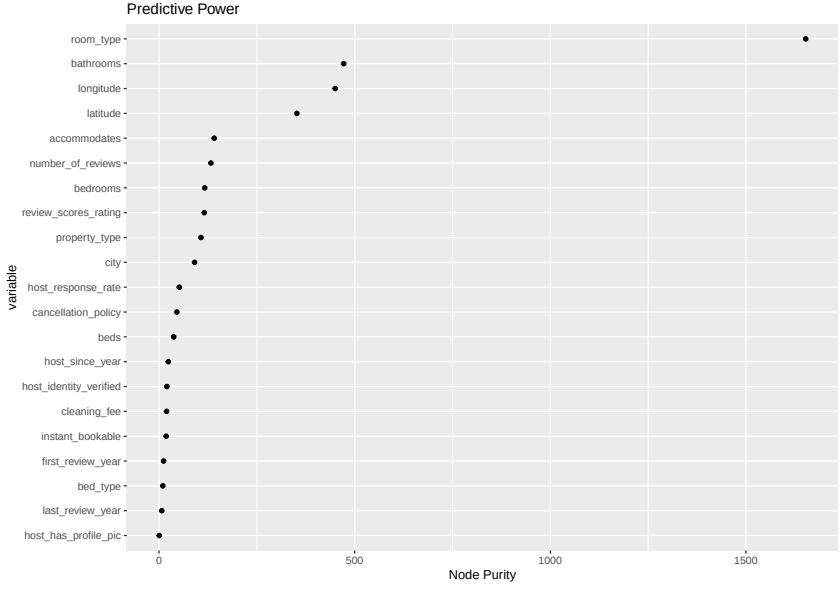
Predictive Power
room_type
bathrooms
longitude
latitude
accommodates
number_of_reviews
bedrooms
review_scores_rating
property_type
city
host_response_rate
cancellation_policy
beds
host_since_year
host_identity_verified
cleaning_fee
instant_bookable
first_review_year
bed_type
last_review_year
host_has_profile_pic
variable
0
500
1000
1500
Node Purity

Random Forests

```
rf_fit <- randomForest(price ~ ., data = training, mtry = s
```

```
## [1] "Test MSE of Random Forest:  0.1299"
```

Random Forests

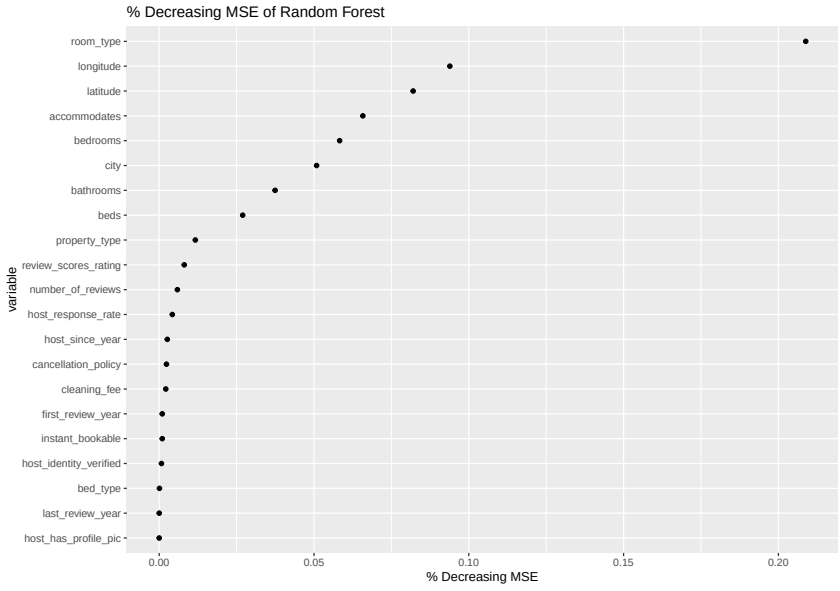

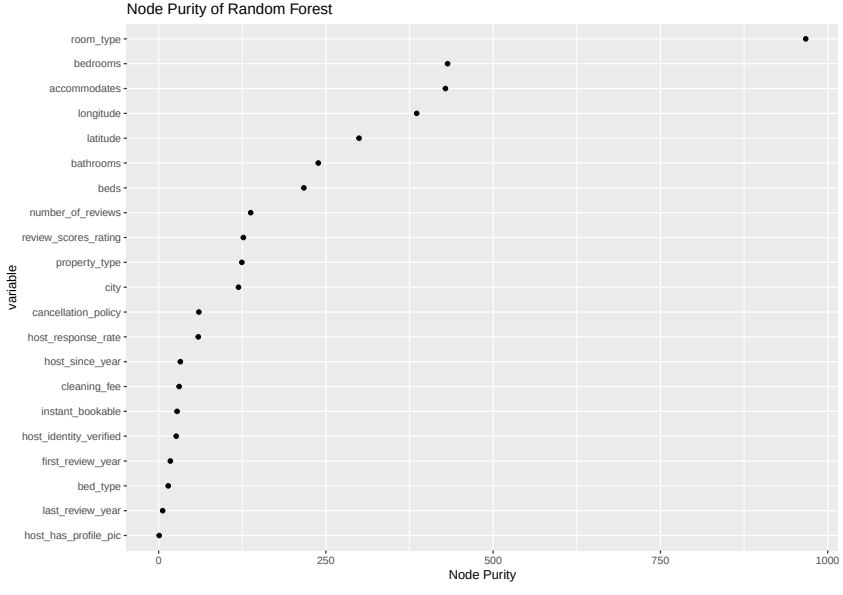

Boosting

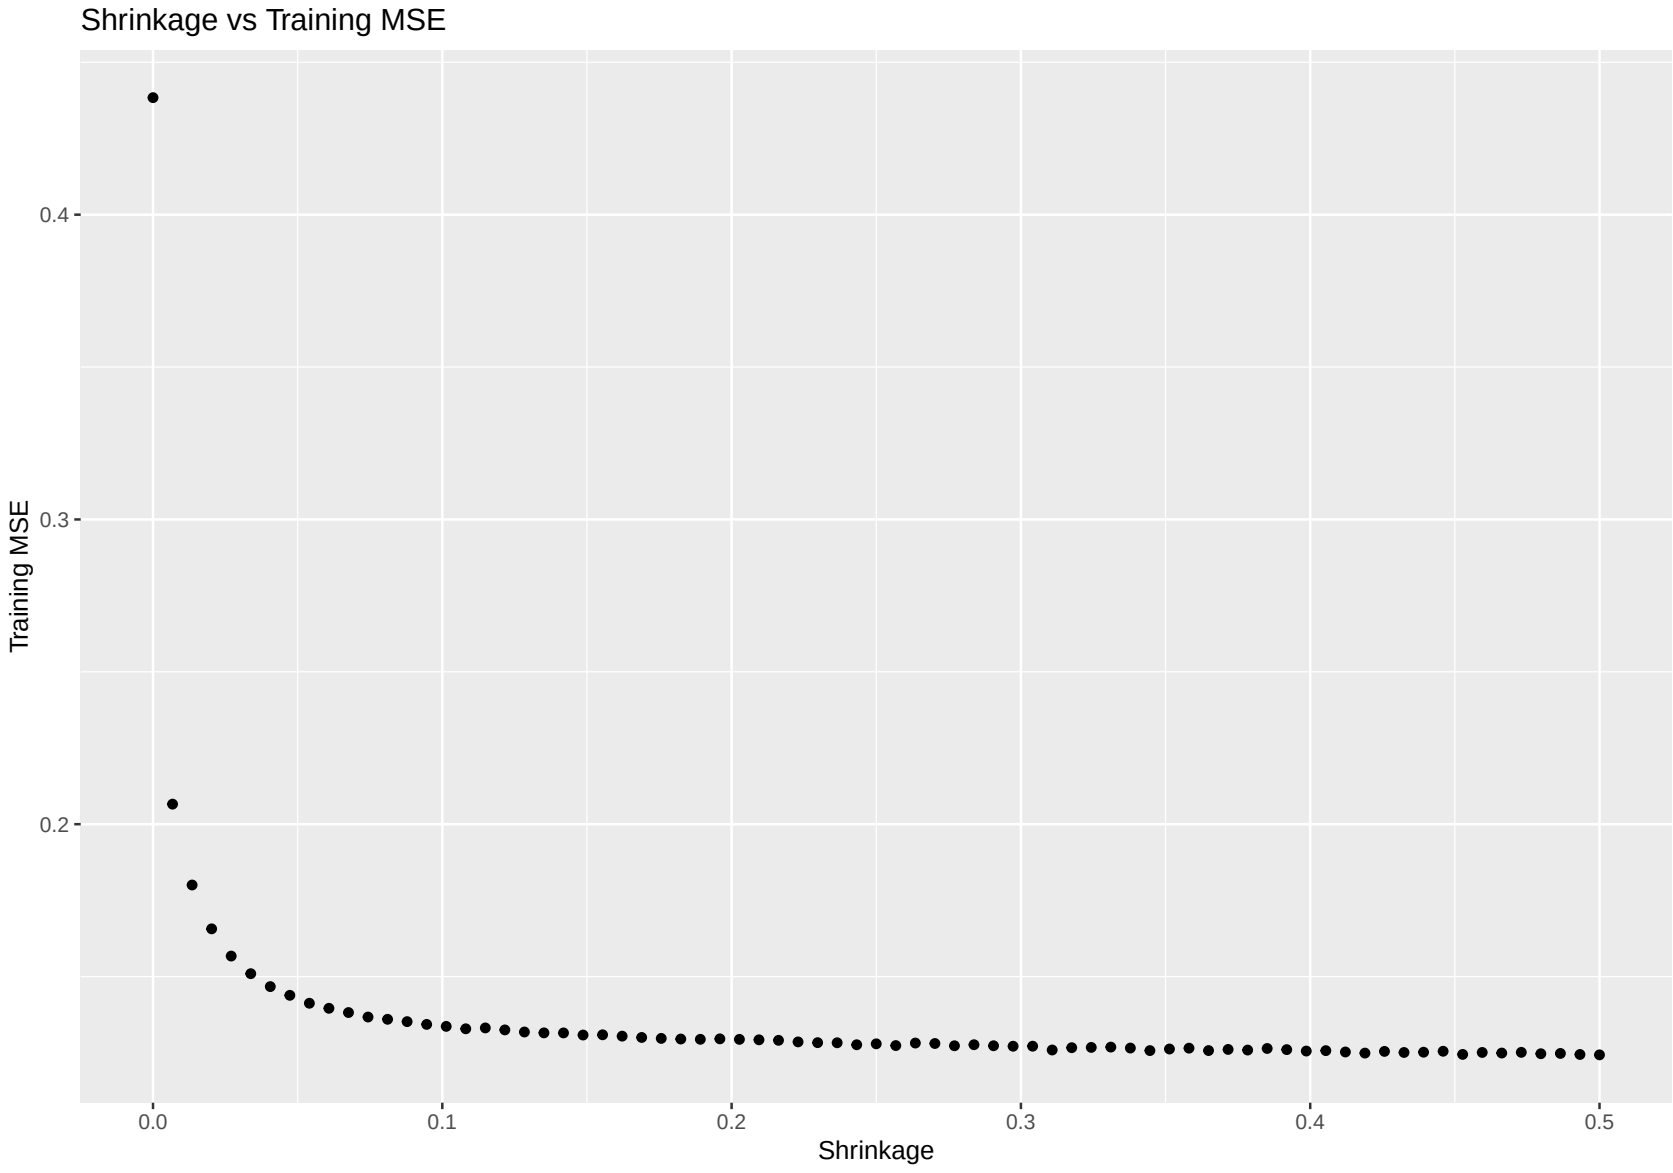
Shrinkage vs Training MSE
Training MSE
0.4
0.3
0.2
0.0
0.1
0.2
0.3
0.4
0.5
Shrinkage

Boosting

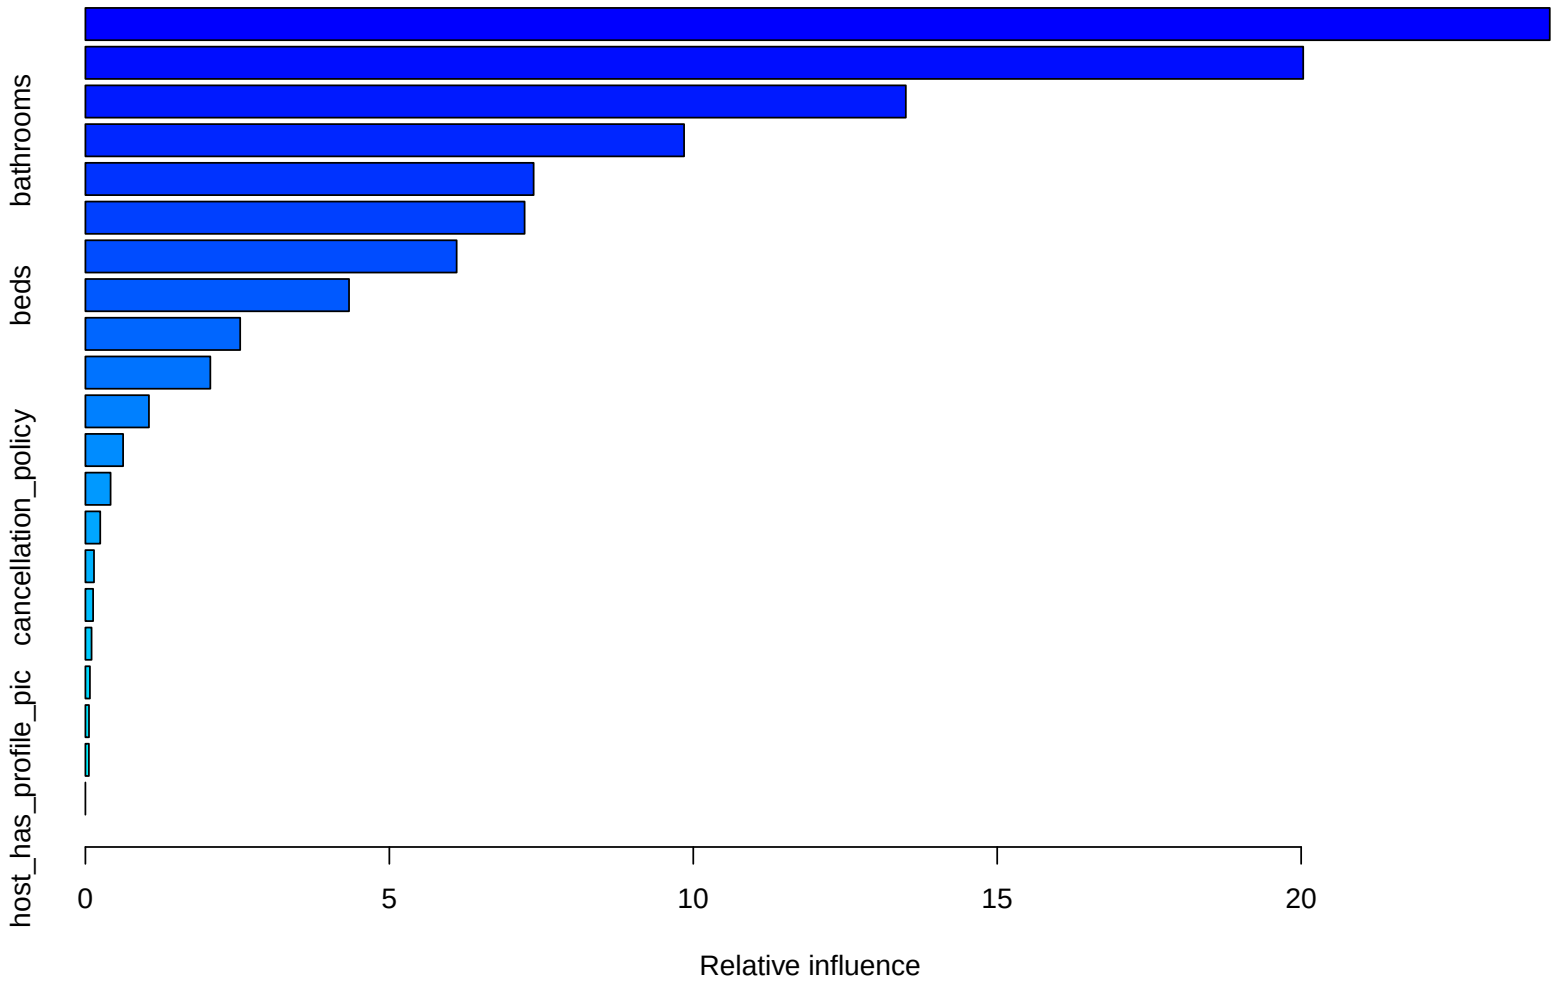


```
##                                    var     rel.in
## property_type                property_type 24.090283
```

MSE Table

Final Model

Going Forward

- Our data has data from multiple cities across the country
- Can we apply this to a certain city and see similar results?
- Is this accurate enough to help AirBnB hosts in selected cities?
  - Using current data, can this model help hosts correctly adjust their rates?

Questions?

References